

# Understanding Operating Systems Seventh Edition

## *Chapter 4* *Processor Management*

# Learning Objectives

After completing this chapter, you should be able to describe:

- The relationship between job scheduling and process scheduling
- The advantages and disadvantages of several process scheduling algorithms
- The goals of process scheduling policies using a single-core CPU

# Learning Objectives (cont'd.)

- The similarities and differences between processes and threads
- The role of internal interrupts and of the interrupt handler

# Overview

- Simple system
  - Single user
  - One processor: busy only when executing the user's job or system software
- Multiprogramming environment: multiple processes competing to be run by a single CPU
  - Requires fair and efficient CPU allocation for each job
- Single processor systems
  - Addressed in this chapter

# Definitions

- Processor (CPU)
  - Performs calculations and executes programs
- Program (job)
  - Inactive unit, e.g., file stored on a disk
  - Unit of work submitted by the user
- Process (task)
  - Active entity
    - Requires resources (processor, special registers, etc.) to perform function
  - Executable program single instance

# Definitions (cont'd.)

- Thread
  - Portion of a process
  - Runs independently
- Multithreading
  - Allows applications to manage a separate process with several threads of control
  - Example: Web browsers
- Multiprogramming
  - Processor allocated to each job or each process for a time period
  - Deallocated at appropriate moment: delicate task

# Definitions (cont'd.)

- Interrupt
  - Call for help
  - Activates higher-priority program
- Context switch
  - Saving job processing information when interrupted
- Completed jobs
  - Finished or terminated
- Single processor
  - May be shared by several jobs (processes)
  - Requires scheduling policy and scheduling algorithm

# About Multi-Core Technologies

- Dual-core, quad-core, or other multi-core CPU
  - More than one processor (core): located on computer chip
- Single chip may contain multiple cores
  - Multi-core engineering
    - Resolves leakage and heat problems
    - Multiple calculations may occur simultaneously
    - More complex Processor Manager handling: discussed in Chapter 6



# Scheduling Submanagers

- Processor Manager: composite of two submanagers
  - Job Scheduler and Process Scheduler: hierarchical scheduling system
- Job Scheduler: higher-level scheduler
  - Job scheduling responsibilities
  - Job initiation based on certain criteria
- Process Scheduler: lower-level scheduler
  - Process scheduling responsibilities
  - Determines execution steps
  - Process scheduling: based on certain criteria

# Scheduling Submanagers (cont'd.)

- Job Scheduler functions
  - Selects incoming job from queue
  - Places in process queue
  - Decides on job initiation criteria
    - Process scheduling algorithm and priority
- Goal
  - Sequence jobs
    - Efficient system resource utilization
  - Balance I/O interaction and computation
  - Keep most system components busy most of time

# Process Scheduler

- Process Scheduler functions
  - Determines job to get CPU resource
    - When and how long
  - Decides interrupt processing
  - Determines queues for job movement during execution
  - Recognizes job conclusion
    - Determines job termination
- Lower-level scheduler in the hierarchy
  - Assigns CPU to execute individual actions: jobs placed on READY queue by the Job Scheduler

# Process Scheduler (cont'd.)

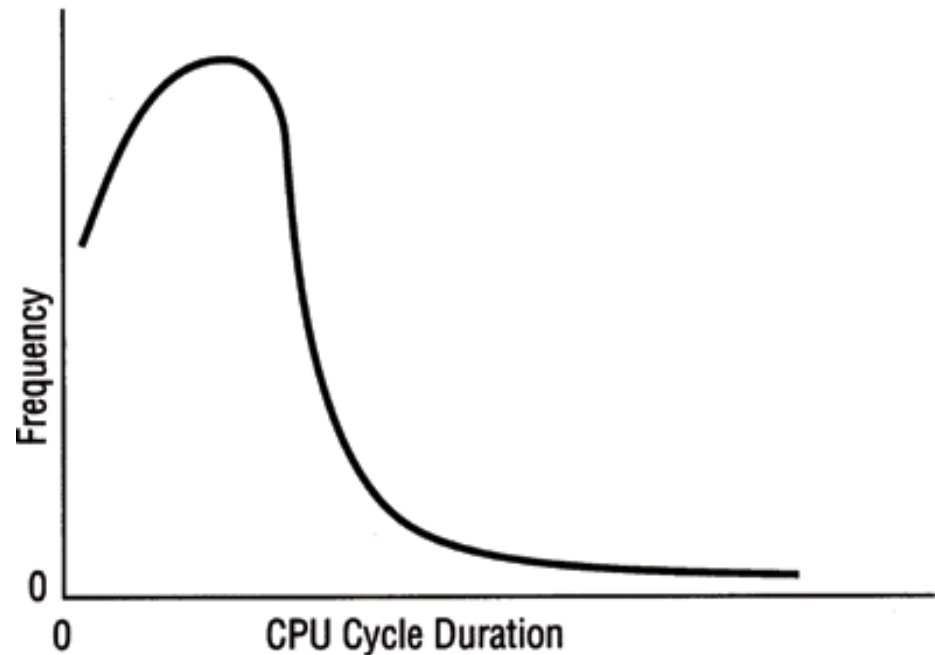
- Exploits common computer program traits
  - Programs alternate between two cycles
    - CPU and I/O cycles
    - Frequency and CPU cycle duration vary
- General tendencies
  - I/O-bound job
    - Many brief CPU cycles and long I/O cycles (printing documents)
  - CPU-bound job
    - Many long CPU cycles and shorter I/O cycles (math calculation)

# Process Scheduler (cont'd.)

**(figure 4.1)**

Distribution of CPU cycle times. This distribution shows a greater number of jobs requesting short CPU cycles (the frequency peaks close to the low end of the CPU cycle axis), and fewer jobs requesting long CPU cycles.

© Cengage Learning 2014



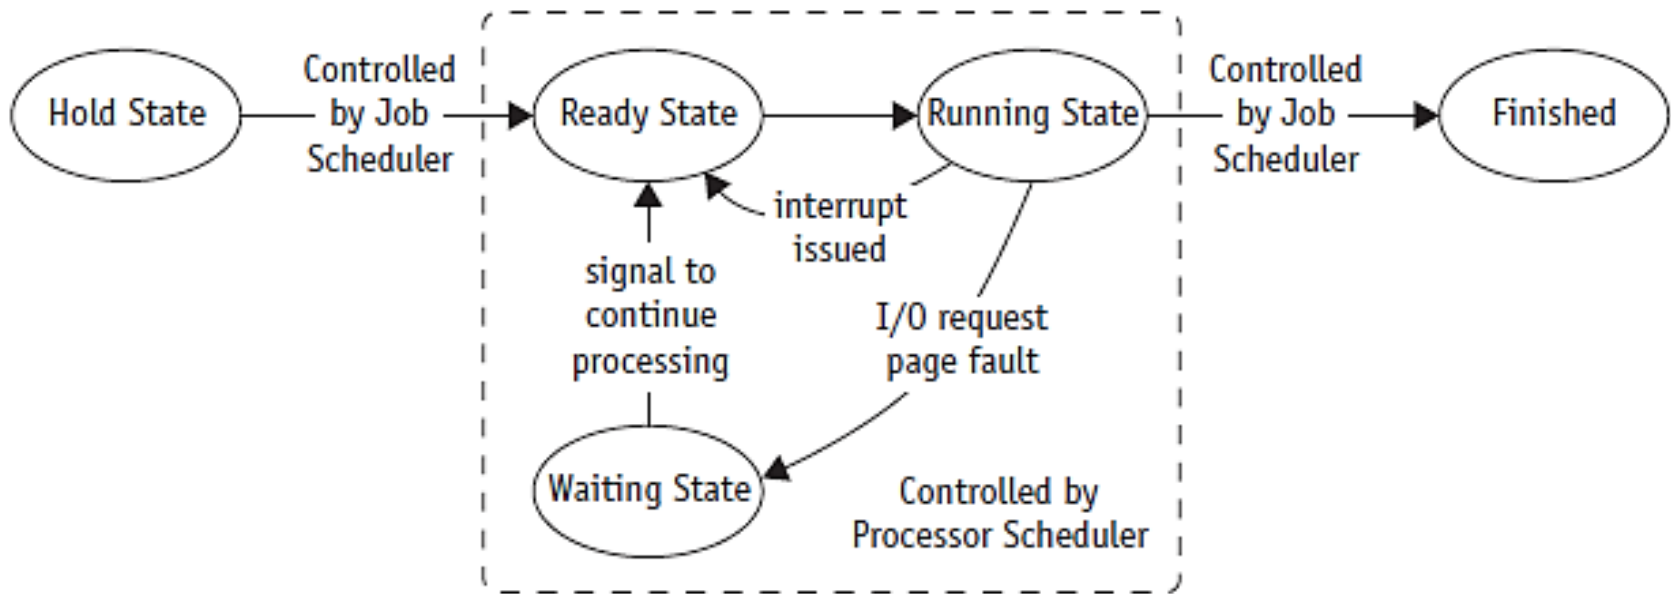
# Process Scheduler (cont'd.)

- Middle-level scheduler: third layer
- Found in highly interactive environments
  - Handles overloading
    - Removes active jobs from memory
    - Reduces degree of multiprogramming
  - Results in faster job completion
- Single-user environment
  - No distinction between job and process scheduling
  - One job active at a time
    - Receives dedicated system resources for job duration

# Job and Process States

- Status changes: as a job or process moves through the system
  - HOLD
  - READY
  - WAITING
  - RUNNING
  - FINISHED
- Referred to as job status or process status, respectively

# Job and Process States (cont'd.)



**(figure 4.2)**

A typical job (or process) changes status as it moves through the system from HOLD to FINISHED.

© Cengage Learning 2014



# Job and Process States (cont'd.)

- User submits job
  - Job accepted
    - Put on HOLD and placed in queue
  - Job state changes from HOLD to READY
    - Indicates job waiting for CPU
  - Job state changes from READY to RUNNING
    - When selected for CPU and processing
  - Job state changes from RUNNING to WAITING
    - Requires unavailable resources: moves back to READY status
  - Job state changes to FINISHED
    - Job completed (successfully or unsuccessfully)

# Job and Process States (cont'd.)

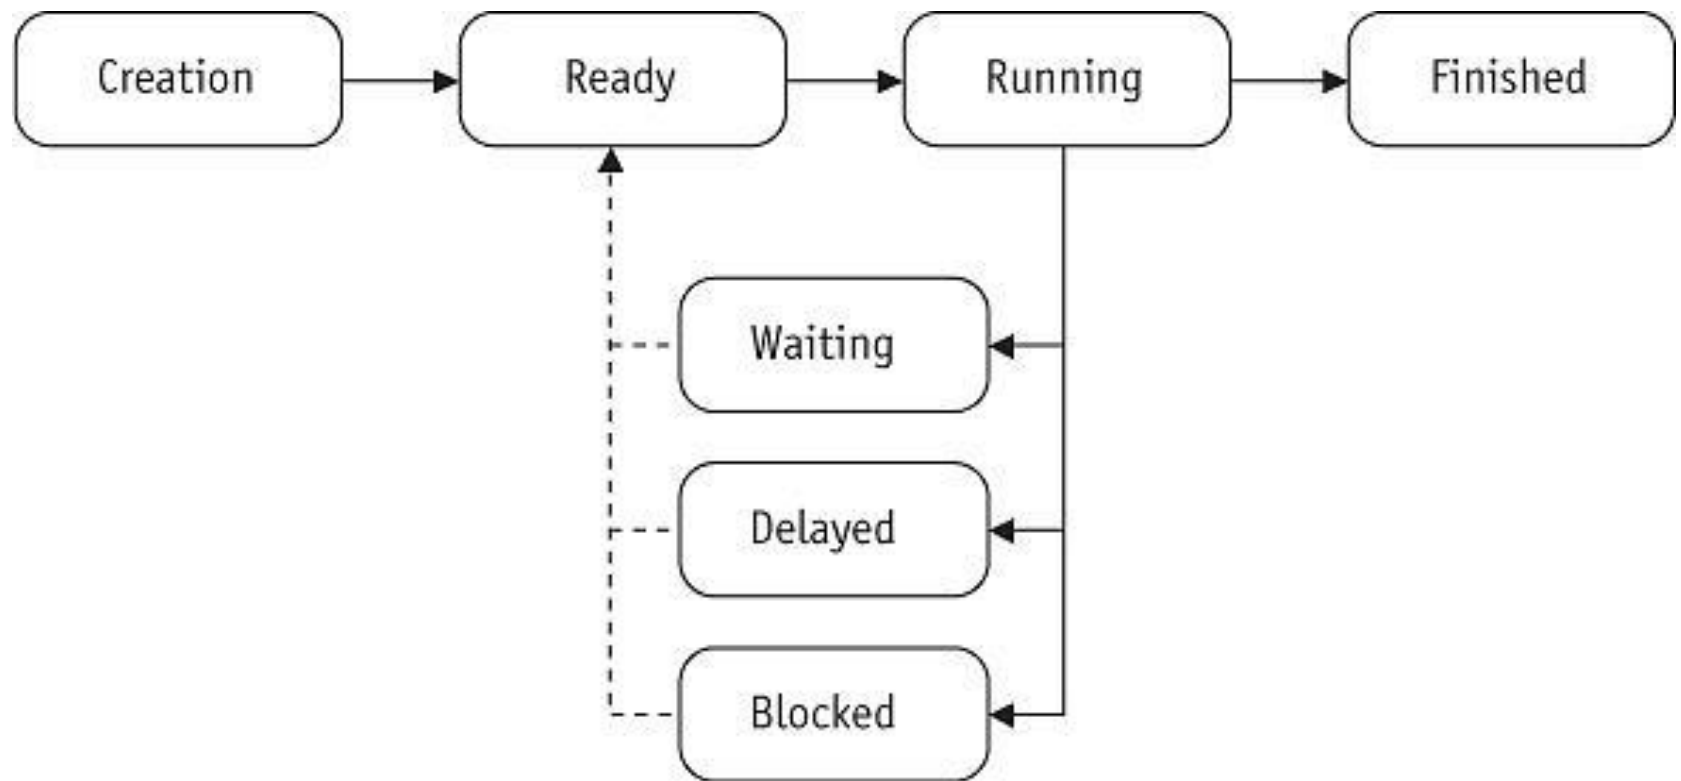
- Job Scheduler or Process Scheduler incurs state transition responsibility
  - HOLD to READY
    - Job Scheduler initiates using predefined policy
  - READY to RUNNING
    - Process Scheduler initiates using predefined algorithm
  - RUNNING back to READY
    - Process Scheduler initiates according to predefined time limit or other criterion
  - RUNNING to WAITING
    - Process Scheduler initiates by instruction in job

# Job and Process Status (cont'd.)

- Job Scheduler or Process Scheduler incurs state transition responsibility (cont'd.)
  - WAITING to READY
    - Process Scheduler initiates by signal from I/O device manager
    - Signal indicates I/O request satisfied; job continues
  - RUNNING to FINISHED
    - Process Scheduler or Job Scheduler initiates upon job completion
    - Satisfactorily or with error

# Thread States

- Five states as a thread moves through the system
  - READY
  - RUNNING
  - WAITING
  - DELAYED
  - BLOCKED
- Thread transitions
  - Application creates a thread: placed in READY queue
  - READY to RUNNING: Process Scheduler assigns it to a processor



**(figure 4.3)**

A typical thread changes states from READY to FINISHED as it moves through the system.

© Cengage Learning 2014

# Thread States (cont'd.)

- Thread transitions
  - Application creates a thread: placed in READY queue
  - READY to RUNNING: Process Scheduler assigns it to a processor
  - RUNNING to WAITING: when dependent on an outside event, e.g., mouse click, or waiting for another thread to finish
  - WAITING to READY: outside event occurs or previous thread finishes

# Thread States (cont'd.)

- Thread transitions (cont'd.)
  - RUNNING to DELAYED: application that delays thread processing by specified amount of time
  - DELAYED to READY: prescribed time elapsed
  - RUNNING to BLOCKED: I/O request issued
  - BLOCKED to RUNNING: I/O completed
  - RUNNING to FINISHED: exit or termination
    - All resources released

# Thread States (cont'd.)

- Operating systems must be able to:
  - Create new threads
  - Set up a thread so it is ready to execute
  - Delay, or put to sleep, threads for a specified amount of time
  - Block, or suspend, threads waiting for I/O to be completed
  - Set threads to a WAIT state until a specific event occurs
  - Schedule threads for execution



# Thread States (cont'd.)

- Operating systems must be able to:
  - Synchronize thread execution using semaphores, events, or conditional variables
  - Terminate a thread and release its resources
- Semaphores are integer variables that are used to solve the critical section problem by using two atomic operations, wait and signal that are used for process synchronization.

# Threads vs. Processes

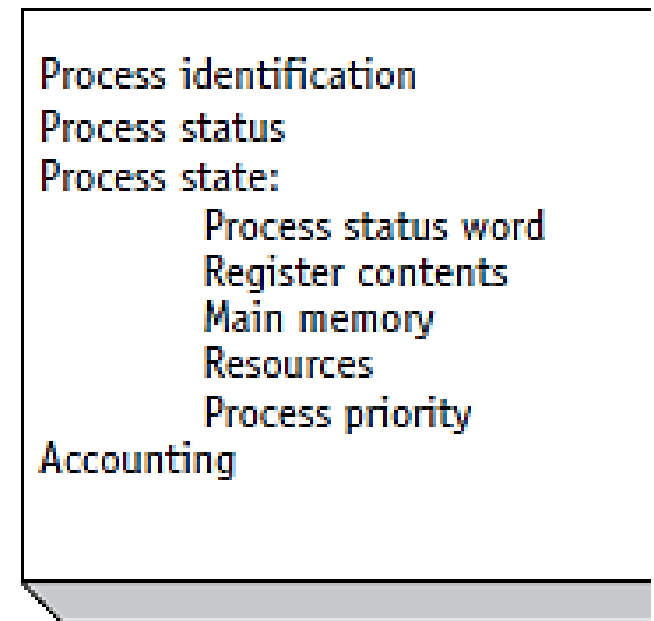
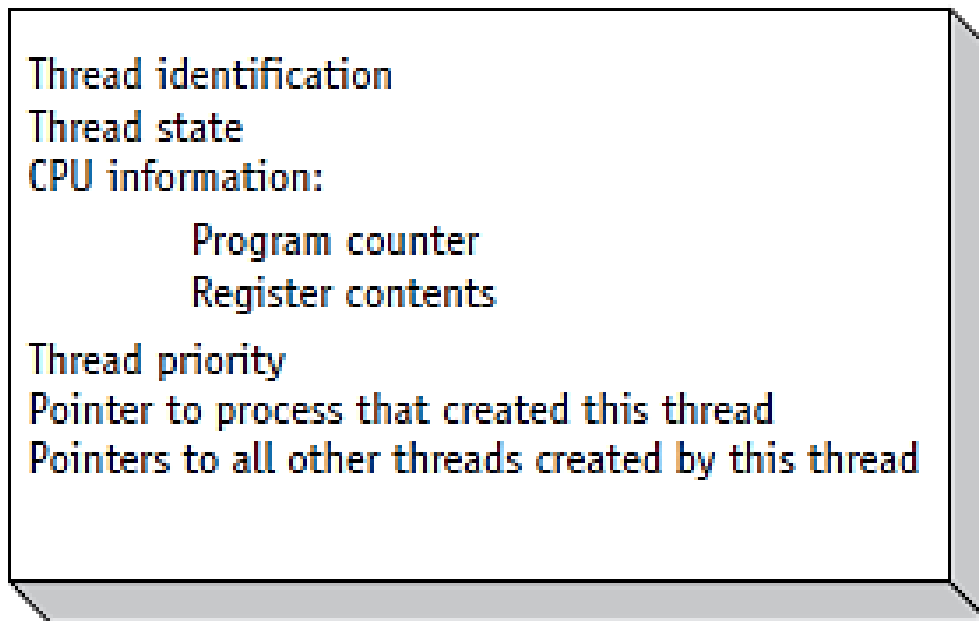
- The program starts out as a text file of programming code,
- The program is compiled or interpreted into binary form,
- The program is loaded into memory,
- The program becomes one or more running processes.
- Processes are typically independent of each other,
- While threads exist as the subset of a process.
- Threads can communicate with each other more easily than processes can,
- But threads are more vulnerable to problems caused by other threads in the same process.

# Processes vs. Threads — Advantages and Disadvantages

PROCESS	THREAD
Processes are heavyweight operations	Threads are lighter weight operations
Each process has its own memory space	Threads use the memory of the process they belong to
Inter-process communication is slow as processes have different memory addresses	Inter-thread communication can be faster than inter-process communication because threads of the same process share memory with the process they belong to
Context switching between processes is more expensive	Context switching between threads of the same process is less expensive
Processes don't share memory with other processes	Threads share memory with other threads of the same process

# Control Blocks

- Process Control Block (PCB): data structure for each process in the system
- Thread Control Block (TCB): data structure for each thread



**(figure 4.4)**

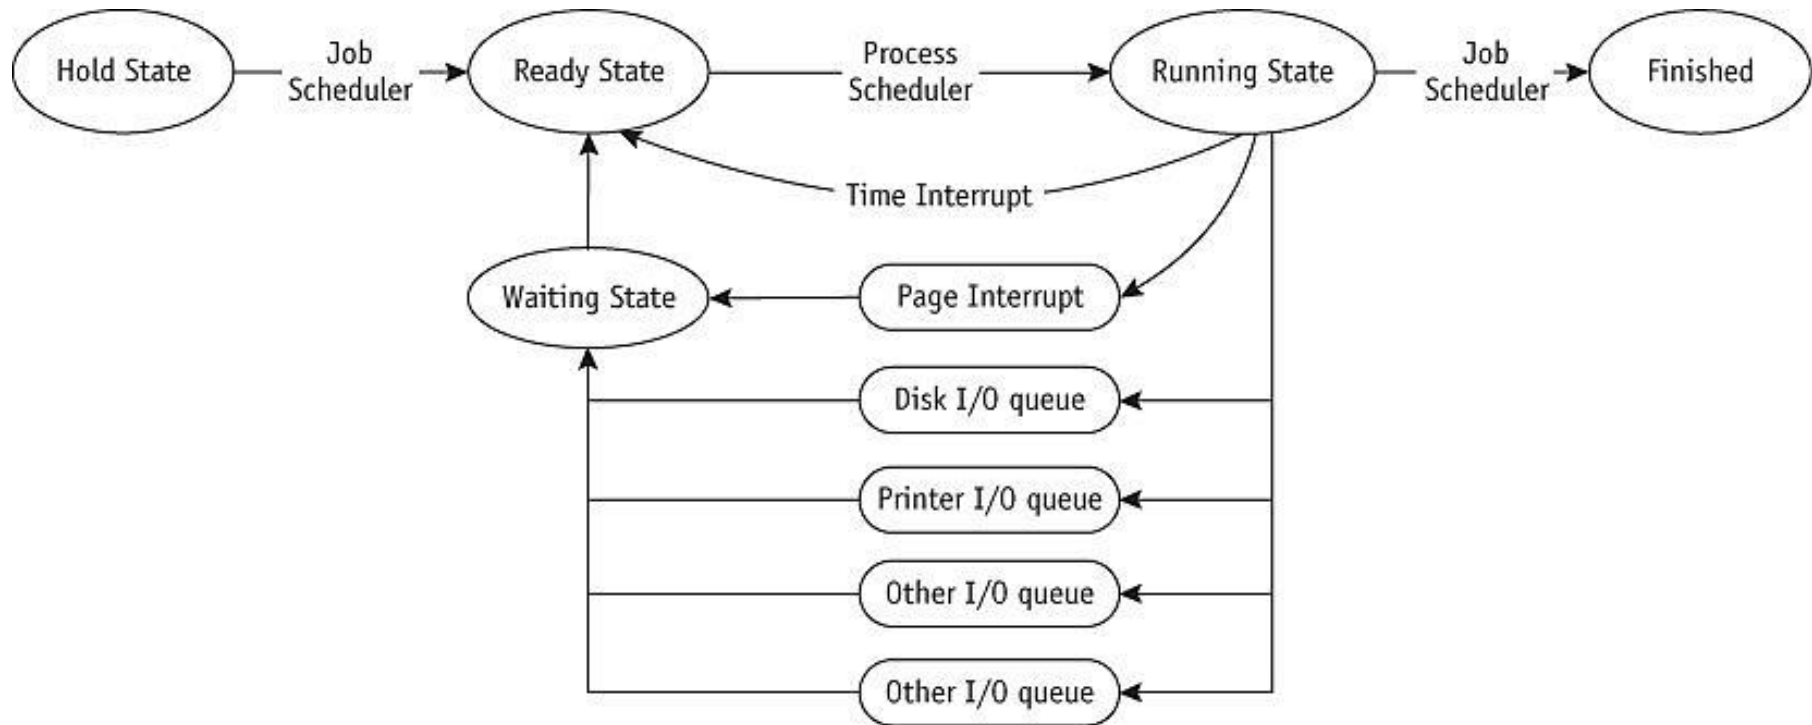
Comparison of a typical Thread Control Block (TCB) vs. a Process Control Block (PCB).

© Cengage Learning 2014

# Control Blocks and Queuing

- Job PCB
  - Created when Job Scheduler accepts job
  - Updated as job executes
  - Queues use PCBs to track jobs
  - Contains all necessary job management processing data
  - PCBs linked to form queues (jobs not linked)
- PCBs or TCBs: requires orderly management of queues
  - Determined by process scheduling policies and algorithms

# Control Blocks and Queuing (cont'd.)



**(figure 4.5)**

Queuing paths from HOLD to FINISHED. The Job and Processor schedulers release the resources when the job leaves the RUNNING state.

© Cengage Learning 2014

# Scheduling Policies

- Multiprogramming environment
  - More jobs than resources at any given time
- Operating system pre-scheduling task
  - Resolve three system limitations
    - Finite number of resources (disk drives, printers, tape drives)
    - Some resources cannot be shared once allocated (printers)
    - Some resources require operator intervention before reassigning



# Scheduling Policies (cont'd.)

- Good process scheduling policy criteria
  - Maximize throughput
    - Run as many jobs as possible in given amount of time
  - Minimize response time
    - Quickly turn around interactive requests
  - Minimize turnaround time
    - Move entire job in and out of system quickly
  - Minimize waiting time
    - Move job out of READY queue quickly

# Scheduling Policies (cont'd.)

- Good process scheduling policy criteria (cont'd.)
  - Maximize CPU efficiency
    - Keep CPU busy 100 percent of time
  - Ensure fairness for all jobs
    - Give every job equal CPU and I/O time
- Final policy criteria decision lies with system designer or administrator

# Scheduling Policies (cont'd.)

- Problem
  - Job claims CPU for very long time before I/O request issued
    - Builds up READY queue and empties I/O queues
    - Creates unacceptable system imbalance
- Corrective measure
  - Interrupt
    - Used by Process Scheduler upon predetermined expiration of time slice
    - Current job activity suspended
    - Reschedules job into READY queue

NOTE: An I/O request is called a natural wait in multiprogramming environments (it allows the processor to be allocated to another job).

# Scheduling Policies (cont'd.)

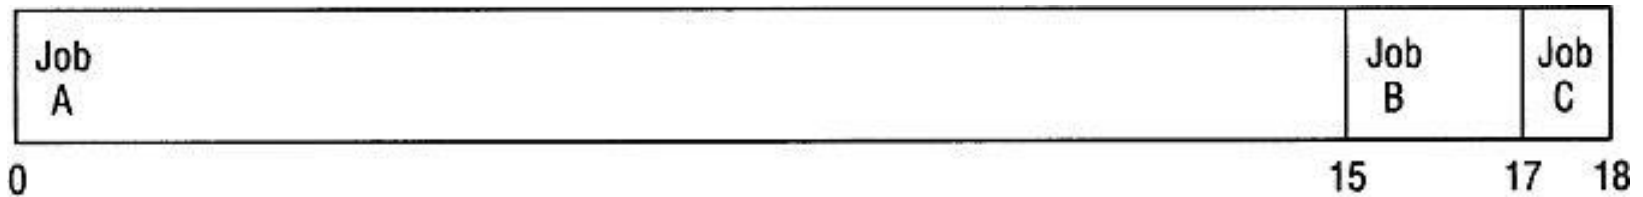
- Types of scheduling policies
  - Preemptive
    - Used in time-sharing environments
    - Interrupts job processing
    - Transfers CPU to another job
  - Nonpreemptive
    - Functions without external interrupts
  - Infinite loops interrupted in both cases

# Scheduling Algorithms

- Based on specific policy
  - Allocates CPU and moves job through the system
- Most systems emphasize fast user response time
- Several algorithms
  - First-come, first-served (FCFS)
  - Shortest job next (SJN)
  - Priority scheduling
  - Shortest remaining time (SRT)
  - Round robin
  - Multiple-level queues
  - Earliest deadline first (EDF)

# First-Come, First-Served

- Nonpreemptive
- Job handled based on arrival time
  - Earlier job arrives, earlier served
- Simple algorithm implementation
  - Uses first-in, first-out (FIFO) queue
- Good for batch systems
- Unacceptable in interactive systems
  - Unpredictable turnaround time
- Disadvantages
  - Average turnaround time varies; seldom minimized

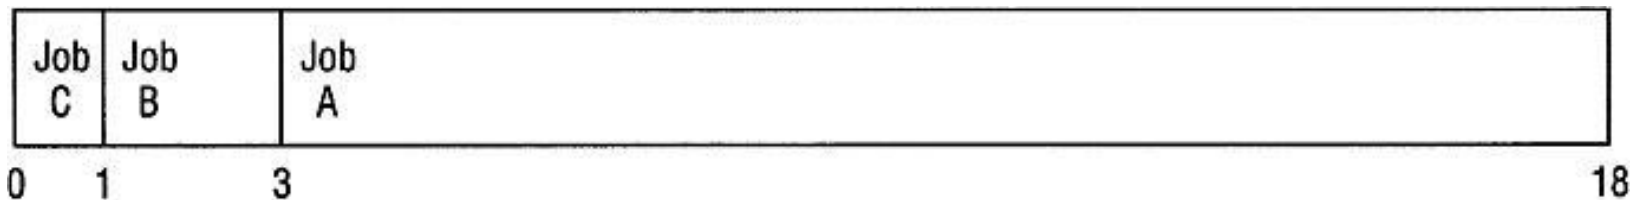


**(figure 4.6)**

Timeline for job sequence A, B, C using the FCFS algorithm.

© Cengage Learning 2014

Note: Average turnaround time: 16.67



**(figure 4.7)**

Timeline for job sequence C, B, A using the FCFS algorithm.

© Cengage Learning 2014

Note: Average turnaround time: 7.3

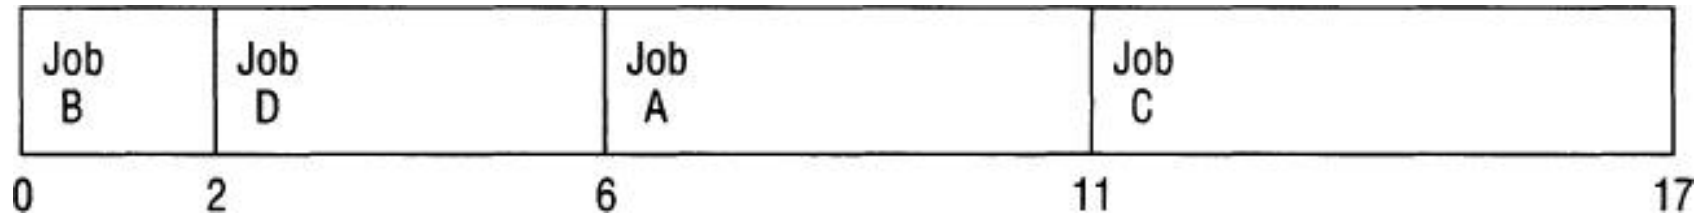
# Shortest Job Next

- Nonpreemptive
- Also known as shortest job first (SJF)
- Job handled based on length of CPU cycle time
- Easy implementation in batch environment
  - CPU time requirement known in advance
- Does not work well in interactive systems
- Optimal algorithm
  - All jobs are available at same time
  - CPU estimates available and accurate



# Shortest Job Next (cont'd.)

Job:	A	B	C	D
CPU cycle:	5	2	6	4



**(figure 4.8)**

Timeline for job sequence B, D, A, C using the SJN algorithm.

© Cengage Learning 2014

# Priority Scheduling

- Nonpreemptive
- Preferential treatment for important jobs
  - Highest priority programs processed first
  - No interrupts until CPU cycles completed or natural wait occurs
- Process Scheduler: manages multiple READY queues
- System administrator or Processor Manager use different methods of assigning priorities

# Priority Scheduling (cont'd.)

- Processor Manager priority assignment methods
  - Memory requirements
    - Jobs requiring large amounts of memory
    - Allocated lower priorities (vice versa)
  - Number and type of peripheral devices
    - Jobs requiring many peripheral devices
    - Allocated lower priorities (vice versa)

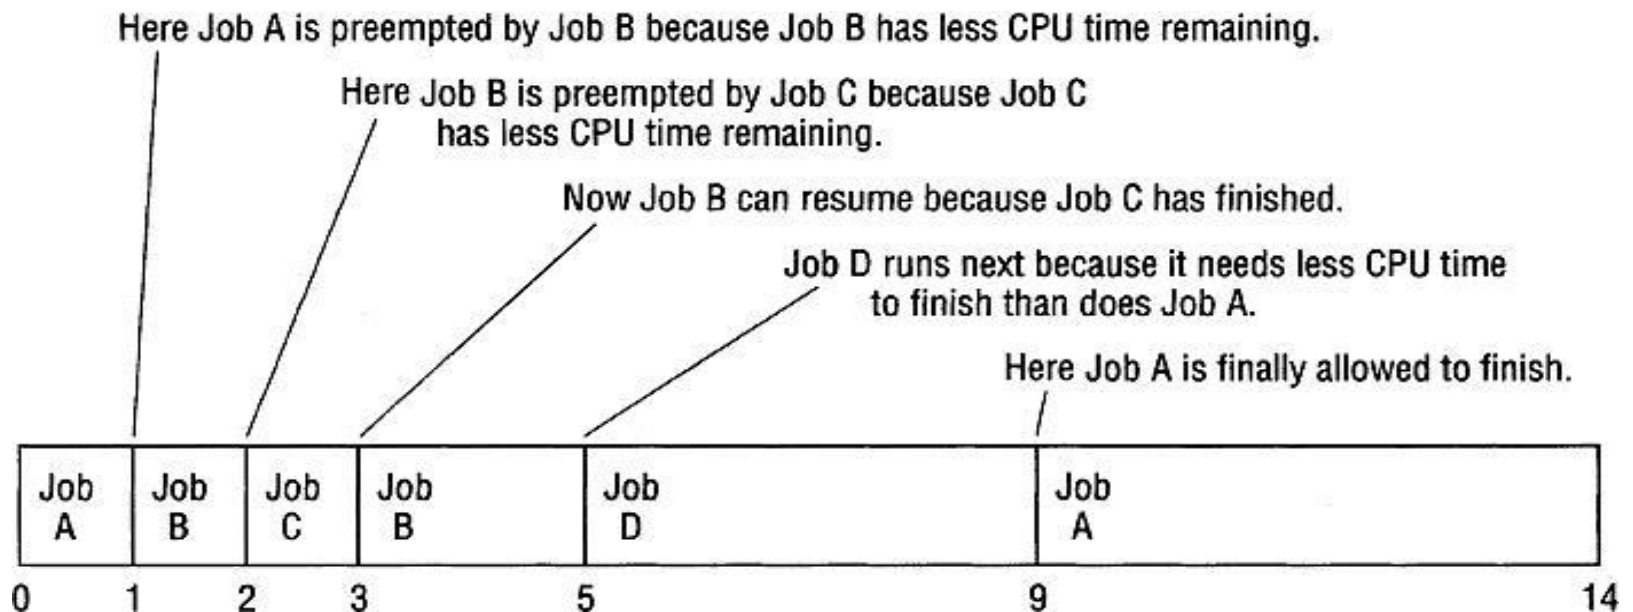
# Priority Scheduling (cont'd.)

- Processor Manager priority assignment methods (cont'd.)
  - Total CPU time
    - Jobs having a long CPU cycle
    - Given lower priorities (vice versa)
  - Amount of time already spent in the system (aging)
    - Total time elapsed since job accepted for processing
    - Increase priority if job in system unusually long time

# Shortest Remaining Time

- Preemptive version of SJN
- Processor allocated to job closest to completion
  - Preemptive if newer job has shorter completion time
- Often used in batch environments
  - Short jobs given priority
- Cannot implement in interactive system
  - Requires advance CPU time knowledge
- Involves more overhead than SJN
  - System monitors CPU time for READY queue jobs
  - Performs context switching

Arrival time:	0	1	2	3
Job:	A	B	C	D
CPU cycle:	6	3	1	4



**(figure 4.9)**

Timeline for job sequence A, B, C, D using the preemptive SRT algorithm. Each job is interrupted after one CPU cycle if another job is waiting with less CPU time remaining.

© Cengage Learning 2014

# Round Robin

- Preemptive
- Used extensively in interactive systems
- Based on predetermined time slice (time quantum)
- Each job assigned time quantum
- Time quantum size
  - Crucial to system performance
  - Varies from 100 ms to 1-2 seconds
- CPU equally shared among all active processes
  - Not monopolized by one job

# Round Robin (cont'd.)

- Job placed on READY queue (FCFS scheme)
- Process Scheduler selects first job
  - Sets timer to time quantum
  - Allocates CPU
- Timer expires
- If job CPU cycle > time quantum
  - Job preempted and placed at end of READY queue
  - Information saved in PCB



# Round Robin (cont'd.)

- If job CPU cycle  $<$  time quantum
  - Job finished: allocated resources released, and job returned to user
  - Interrupted by I/O request: information saved in PCB and linked to I/O queue
- Once I/O request satisfied
  - Job returns to end of READY queue and awaits CPU

# Round Robin (cont'd.)

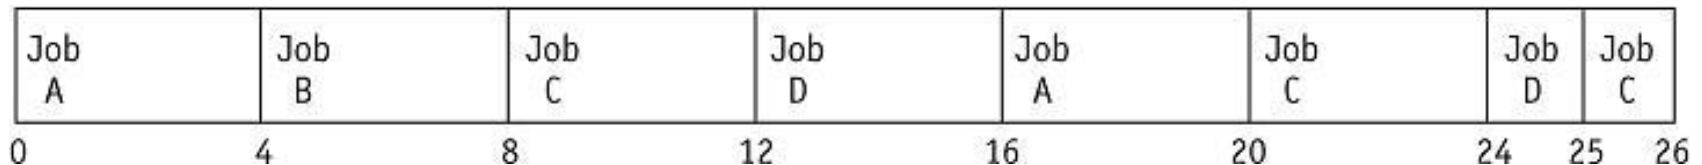
The example in Figure 4.11 illustrates a Round Robin algorithm with a time slice of 4 milliseconds (I/O requests are ignored):

Arrival time:	0	1	2	3	Job:	A	B	C	D
Job:	A	B	C	D	Completion Time minus Arrival Time	20-0	8-1	26-2	25-3
CPU cycle:	8	4	9	5	Turnaround:	20	7	24	22

So the average turnaround time is:

$$\frac{20 + 7 + 24 + 22}{4} = 18.25$$

The turnaround time is the completion time minus the arrival time:



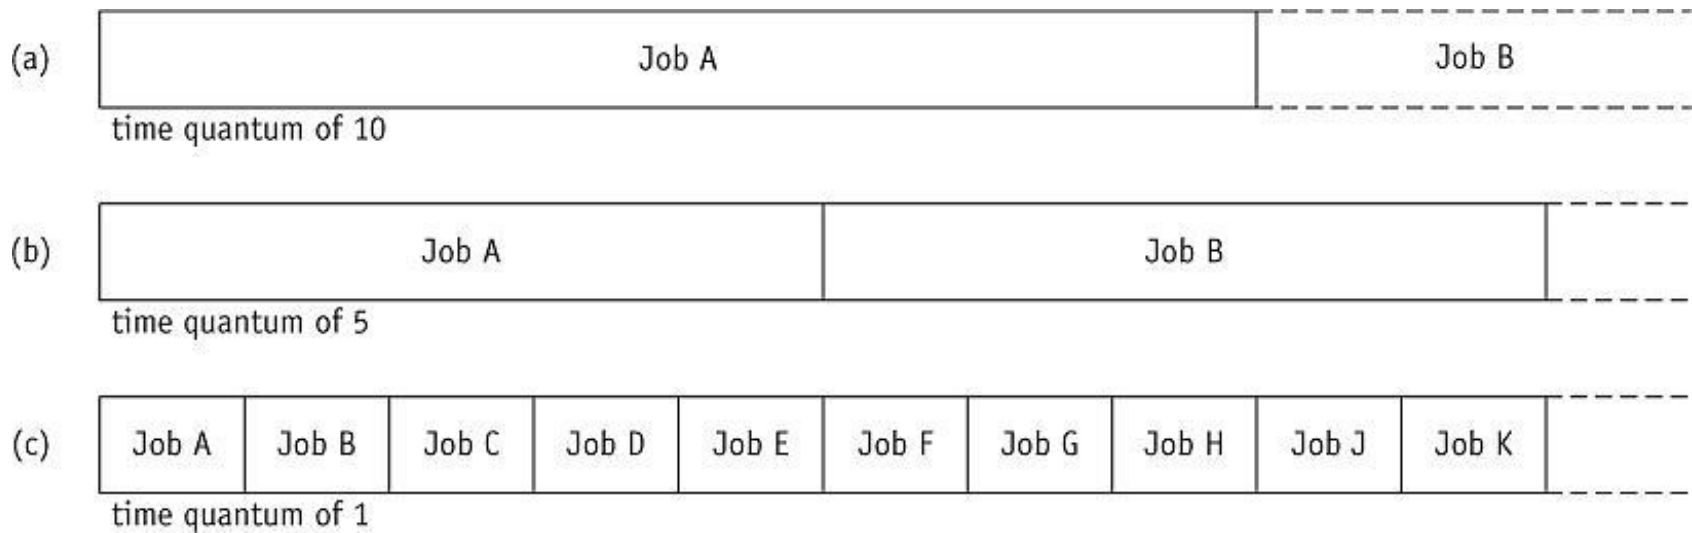
**(figure 4.11)**

Timeline for job sequence A, B, C, D using the preemptive round robin algorithm

# Round Robin (cont'd.)

- Efficiency
  - Depends on time quantum size
    - In relation to average CPU cycle
- Quantum too large (larger than most CPU cycles)
  - Algorithm reduces to FCFS scheme
- Quantum too small
  - Context switching occurs
    - Job execution slows down
    - Overhead dramatically increased

# Round Robin (cont'd.)



**(figure 4.12)**

Context switches for three different time quantum. In (a), Job A (which requires only 8 cycles to run to completion) finishes before the time quantum of 10 expires. In (b) and (c), the time quantum expires first, interrupting the jobs.

© Cengage Learning 2014

# Round Robin (cont'd.)

- Best quantum time size
  - Depends on system
    - Interactive: response time key factor
    - Archival system: turnaround time key factor
  - General rules of thumb
    - Long enough for 80% of CPU cycles to complete
    - At least 100 times longer than context switch time requirement

# Multiple-Level Queues

- Works in conjunction with several other schemes
- Works well in systems with jobs grouped by common characteristic
  - Priority-based
    - Different queues for each priority level
  - CPU-bound jobs in one queue and I/O-bound jobs in another queue
  - Hybrid environment
    - Batch jobs in background queue
    - Interactive jobs in foreground queue
- Scheduling policy based on predetermined scheme
- Four primary methods of moving jobs

# Multiple-Level Queues (cont.)

- Multiple-level queues raise some interesting questions:
  - Is the processor allocated to the jobs in the first queue until it is empty before moving to the next queue, or does it travel from queue to queue until the last job on the last queue has been served? And then go back to serve the first job on the first queue? Or something in between?
  - Is this fair to those who have earned, or paid for, a higher priority?
  - Is it fair to those in a low-priority queue?
  - If the processor is allocated to the jobs on the first queue and it never empties out, when will the jobs in the last queues be served?
  - Can the jobs in the last queues get “time off for good behavior” and eventually move to better queues?

# Multiple-Level Queues (cont.)

- The answers depend on the policy used by the system to service the queues.
- Four primary methods to the movement:
  - Not allowing movement between queues,
  - Moving jobs from queue to queue,
  - Moving jobs from queue to queue and increasing the time quantum for lower queues,
  - Giving special treatment to jobs that have been in the system for a long time (aging).



# Case 1: No Movement Between Queues

- Simple
- Rewards high-priority jobs
  - Processor allocated using FCFS
- Processor allocated to lower-priority jobs
  - Only when high-priority queues empty
- Good environment
  - Few high-priority jobs
  - Spend more time with low-priority jobs

# Case 2: Movement Between Queues

- Processor adjusts priorities assigned to each job
- High-priority jobs
  - Initial priority favorable
    - Treated like all other jobs afterwards
- Quantum interrupt
  - Job preempted
    - Moved to next lower queue
    - May have priority increased
- Good environment
  - Jobs handled by cycle characteristics (CPU or I/O)
  - Interactive systems

## Case 3: Variable Time Quantum Per Queue

- Case 2 variation: movement between queues
- Each queue given time quantum size
  - Size twice as long as previous queue
- Fast turnaround for CPU-bound jobs
- CPU-bound jobs execute longer and given longer time periods
  - Improves chance of finishing faster

## Case 4: Aging

- Ensures lower-level queue jobs eventually complete execution
- System keeps track of job wait time
- If too “old”:
  - System moves job to next highest queue
  - Continues until old job reaches top queue
  - May drastically move old job to highest queue
- Advantage
  - Guards against indefinite postponement
    - Major problem: discussed further in Chapter 5

# Earliest Deadline First (EDF)

- Also known as Dynamic priority algorithm
- Preemptive
- Addresses critical processing requirements of real-time systems: deadlines
- Job priorities can be adjusted while moving through the system
- Primary goal:
  - Process all jobs in order most likely to allow each to run to completion before reaching their respective deadlines

# Earliest Deadline First (cont'd.)

- Initial job priority: inversely proportional to its absolute deadline
  - Jobs with same deadlines: another scheme applied
- Priority can change as more important jobs enter the system
- Problems
  - Missed deadlines: total time required for all jobs greater than allocated time until final deadline
  - Impossible to predict job throughput: changing priority nature
  - High overhead: continual evaluation of deadlines

# Managing Interrupts

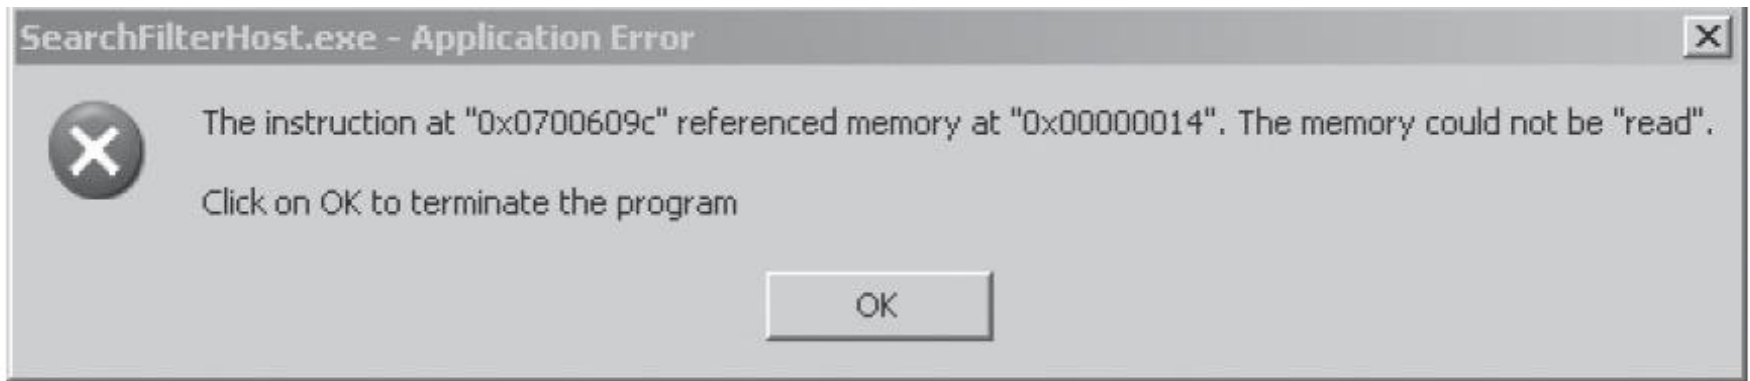
- Interrupt types
  - Page interrupt (memory manager)
    - Accommodate job requests
  - Time quantum expiration interrupt
  - I/O interrupt
    - Result from READ or WRITE command issuance
  - Internal interrupt (synchronous interrupt)
    - Result from arithmetic operation or job instruction
  - Illegal arithmetic operation interrupt
    - Dividing by zero; bad floating-point operation, fixed-point operations resulting in overflow or underflow

# Managing Interrupts (cont'd.)

- Interrupt types (cont'd.)
  - Illegal job instruction interrupt
    - Attempting protected storage access, operating on invalid data, etc.
    - Attempts to use an undefined operation code
    - Unauthorized attempts to make system changes, such as trying to change the size of the time quantum
- Interrupt handler
  - Control program
    - Handles interruption event sequence



# Managing Interrupts (cont'd.)



**(figure 4.14)**

Sample Windows screen showing that the interrupt handler has stepped in after an invalid operation.

# Managing Interrupts (cont'd.)

- Nonrecoverable error detected by operating system
  - Interrupt handler sequence
    1. Interrupt type described and stored
    2. Interrupted process state saved
    3. Interrupt processed
    4. Processor resumes normal operation

# Conclusion

- Processor Manager allocates CPU among all users
- Job scheduling
  - Based on characteristics
- Process and thread scheduling
  - Instant-by-instant allocation of CPU
- Interrupt handler: generates and resolves interrupts
- Each scheduling algorithm is unique
  - Characteristics, objectives, and applications
- System designer selects best policy and algorithm
  - After careful strengths and weaknesses evaluation

Algorithm	Policy Type	Disadvantages	Advantages
First Come, First Served	Nonpreemptive	Unpredictable turnaround times; has an element of chance	Easy to implement
Shortest Job Next	Nonpreemptive	Indefinite postponement of some jobs; requires execution times in advance	Minimizes average waiting time
Priority Scheduling	Nonpreemptive	Indefinite postponement of some jobs	Ensures fast completion of important jobs
Shortest Remaining Time	Preemptive	Overhead incurred by context switching	Ensures fast completion of short jobs
Round Robin	Preemptive	Requires selection of good time quantum	Provides reasonable response times to interactive users; provides fair CPU allocation

**(table 4.3)**

Comparison of the scheduling algorithms discussed in this chapter.

© Cengage Learning 2014

Algorithm	Policy Type	Disadvantages	Advantages
Multiple-Level Queues	Preemptive/ Nonpreemptive	Overhead incurred by monitoring queues	Flexible scheme; allows aging or other queue movement to counteract indefinite postponement; is fair to CPU-bound jobs
Earliest Deadline First	Preemptive	Overhead required to monitor dynamic deadlines	Attempts timely completion of jobs

**(table 4.3) (cont'd.)**

Comparison of the scheduling algorithms discussed in this chapter.

© Cengage Learning 2014